

OS2 CS3523

Name: Swapnil Bag

Roll no: ES22BTECH11034

Aim: To solve the Readers-Writers problem (writer preference) and Fair Readers-Writers problem using Semaphores and thread library of C++

Low level design

- Firstly the input (nr , nw, kr, kw, randRemTime, randCSTime) are accepted using a command line argument
- A structure `_th_param_` is created to store thread parameters like id , kr, kw, randRemTime, randCSTime
- Reader and writer threads are created in the main function and the respective parameters are passed to the reader and writer threads

- **Writer Preference solution**

- 1) Count of readers and writers is initialized to 0. 4 semaphores `r_lock`(reader lock), `w_lock`(writer lock), `r_try`(reader trylock), resource (resource lock) and a `print_l` (to synchronize printing to file).
- 2) In the reader function, each thread makes `kr` requests to read. First it enters and tries to acquire the `r_try` lock and the `r_lock` (so that other readers do not enter simultaneously). The count of readers is incremented and if there is a single reader the resources is locked ,it releases the `r_lock` and the `r_try` lock enters the critical section where it sleeps for a random time and then exits the critical section by invoking `r_lock` and if there are no more readers it signals the writers.
- 3) The writer threads makes `kw` requests to read, it acquires the `w_lock`, increments count of writers, if there is a single writer you lock the `r_try` semaphore so that no reader can enter, release `w_lock` and acquire resource. In the critical section after finishing it releases resource. While exiting it updates the count of writers and releases the `r_try`
- 4) The main purpose of the `r_try` lock is that whenever a reader wants to access the CS, each reader should individually acquire the `r_try` and release it, but in case of writers if a single writer acquires it then an incoming stream of writers can be serviced (each writer need not individually acquire it). This ensures writers preference but readers can be starved

- **Fair Reader writers solution**

- 1) Count of readers and writers is initialized to 0. 4 semaphores `r_lock`(reader lock), resource (resource lock) and a `print_l` (to synchronize printing to file) ad a `serv_q` lock that in a way ensures fairness among the threads
- 2) In the reader function, each thread makes `kr` requests to read. First it enters and tries to acquire the `serv_q` lock and the `r_lock` (so that other readers do not enter simultaneously). The count of readers is incremented and if there is a single reader the resources is locked ,it releases the `r_lock` and the `serv_q`, lock enters the critical section where it sleeps for a random time and then exits the critical section by invoking `r_lock` and if there are no more readers it signals the writers.
- 3) The writer threads makes `kw` requests to read, it acquires the `serv_q`, and then acquires the resource so that no other thread can enter. While exiting it releases the resource lock.

- 4) The main purpose of the serv_q lock is to ensure fairness among the reader and writer threads. In a way each reader or writer must invoke this and release it individually, so there's an equal chance of both the threads to acquire which in turn ensures fairness.

- **Output** : For each solution The logs is printed in log.txt and Average_time.txt. In the log file the time stamp is provided in hh:mm:ss:microsecond format for better precision, and In the average time file, the average and worst case time from request to entry is printed

Analysis:

We can see that in the writer preference code the average waiting time for reader threads is usually greater than the average waiting time of writer thread , which holds roughly in the worst case also. This worst case time is mainly contributed due to multiple lock contentions. In the log file we can observe that initially readers request for CS however over time the writer threads execute in one chunk as it acquired the lock for themselves, thus resulting in a preference for writers. When almost all writers finish the readers complete

In the fair code The average as well as the worst case time is nearly equal for both type of threads. This is possibly because of the simple lock usage and lock acquiring sequence for both type of threads. Furthermore from the log file we can infer that reader and writer execution is evenly interspersed indicating both type of threads get a fair chance to execute



